

Reading and Importing Data

```
In [1]: #importing csv file
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import scale
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import classification_report
from sklearn.metrics import classification_report, confusion_matrix
data = pd.read_csv("digit_svm.csv")
```

Data understanding and exploration¶

```
In [2]: data.head()
```

```
Out[2]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780
0	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
2	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
3	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0

5 rows × 785 columns

```
In [3]: data.shape
```

```
Out[3]: (42000, 785)
```

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 42000 entries, 0 to 41999
Columns: 785 entries, label to pixel783
dtypes: int64(785)
memory usage: 251.5 MB
```

```
In [5]: data.describe()
```

```
Out[5]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780
count	42000.000000	42000.0	42000.0	42000.0	42000.0	42000.0	42000.0	42000.0	42000.0	42000.0	...	42000.000000	42000.000000	42000.000000	42000.000000	42000.000000	42000.000000	42000.000000
mean	4.456643	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.219286	0.117095	0.059893	0.059893	0.059893	0.059893	0.059893
std	2.887730	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	6.312890	4.633819	3.274178	3.274178	3.274178	3.274178	3.274178
min	0.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	2.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	4.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	7.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	9.000000	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	254.000000	254.000000	254.000000	254.000000	254.000000	254.000000	254.000000

8 rows × 785 columns

```
In [6]: data.isnull().sum()
```

```
Out[6]: label      0
pixel0      0
```

```

pixel1      0
pixel2      0
pixel3      0
...
pixel779    0
pixel780    0
pixel781    0
pixel782    0
pixel783    0
Length: 785, dtype: int64

```


Viualisation

```
In [7]: sns.countplot(data["label"])
```

C:\Users\KIIT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[7]: <AxesSubplot:xlabel='label', ylabel='count'>
```




```
In [8]: plt.plot(figsize = (16,10))
g = sns.countplot( data["label"], palette = 'icefire')
plt.title('NUmber of digit classes')
data.label.astype('category').value_counts()
```

C:\Users\KIIT\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

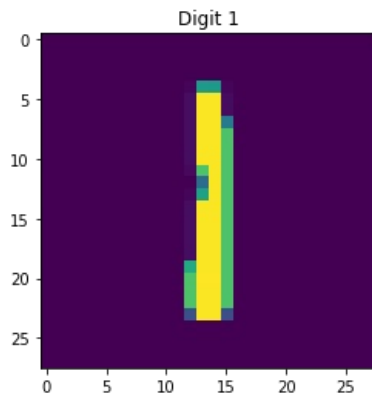
```
warnings.warn(
```

```
Out[8]: 1    4684
       7    4401
       3    4351
       9    4188
       2    4177
       6    4137
       0    4132
       4    4072
       8    4063
       5    3795
Name: label, dtype: int64
```




```
In [9]: one = data.iloc[2, 1:]
one = one.values.reshape(28,28)
plt.imshow(one)
plt.title("Digit 1")
```

Out[9]: Text(0.5, 1.0, 'Digit 1')




```
In [10]: four = data.iloc[3, 1:]
four = four.values.reshape(28,28)
plt.imshow(four)
plt.title("Digit 4")
```

Out[10]: Text(0.5, 1.0, 'Digit 4')



```
In [11]: nine = data.iloc[11, 1:]
nine = nine.values.reshape(28,28)
plt.imshow(nine)
plt.title("Digit 9")
```


Out[11]: Text(0.5, 1.0, 'Digit 9')



Heatmap

```
In [12]: plt.figure(figsize=(15,15))
sns.heatmap(data=data.corr(),annot=False)
```

```
Out[12]: <AxesSubplot:>
```



```
In [13]: data.corr()
```

```
Out[13]:
```

	label	pixel0	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	...	pixel774	pixel775	pixel776	pixel777	pixel778	pixel779
label	1.000000	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	0.033424	0.025050	0.019558	0.014490	0.009790	0.006075
pixel0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
pixel1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
pixel2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
pixel3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
...
pixel779	0.006075	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	-0.000240	-0.000174	-0.000124	0.236633	0.905835	1.000000

pixel780	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN
pixel781	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN
pixel782	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN
pixel783	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN

785 rows × 785 columns

Splitting and scalling

```
In [14]: # splitting into X and y
X = data.drop("label", axis = 1)
y = data['label']
X_scaled = scale(X)
from sklearn.model_selection import train_test_split
# train test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, train_size=0.2, test_size = 0.8, random_state = 1)
```

```
In [15]: print('X_train shape:', X_train.shape)
print('y_train shape:', y_train.shape)
print('X_test shape:', X_test.shape)
print('y_test shape:', y_test.shape)
```

```
X_train shape: (8400, 784)
y_train shape: (8400,)
X_test shape: (33600, 784)
y_test shape: (33600,)
```

Model Building

```
In [16]: from sklearn import svm
from sklearn import metrics
from sklearn.svm import SVC

model_linear = SVC(kernel='linear')
model_linear.fit(X_train, y_train)
y_pred = model_linear.predict(X_test)

# accuracy
print("accuracy:", metrics.accuracy_score(y_true=y_test, y_pred=y_pred), "\n")

# cm
print(metrics.confusion_matrix(y_true=y_test, y_pred=y_pred))
```

accuracy: 0.913125

```
[[3188   0  10   5  11  20  32   3  15   1]
 [   0 3677  14  11   5   7   4   8  30   4]
 [   36   29 3027  54  55  10  30  42  48  12]
 [   13   12 104 3051   9 181   5  21  54  25]
 [    8   14   33   2 3057   4  25  31   6 110]
 [   30   23   29 136  44 2622  44  12  72  27]
 [   26   11   44   4  28   33 3113   0  18   0]
 [    7   24   36  19  59   9   2 3210   4 134]
 [   13   46   50 120  21 110   30  18 2843  21]
 [   19   17   21  22 172   20   4  161  26 2893]]
```

```
In [17]: scores=metrics.classification_report(y_test, y_pred, labels=[0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
print(scores)
```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	3285
1	0.95	0.98	0.97	3760
2	0.90	0.91	0.90	3343
3	0.89	0.88	0.88	3475
4	0.88	0.93	0.91	3290
5	0.87	0.86	0.87	3039

6	0.95	0.95	0.95	3277
7	0.92	0.92	0.92	3504
8	0.91	0.87	0.89	3272
9	0.90	0.86	0.88	3355
accuracy			0.91	33600
macro avg	0.91	0.91	0.91	33600
weighted avg	0.91	0.91	0.91	33600

Non-Linear SVM

```
In [18]: # rbf kernel with other hyperparameters kept to default
svm_rbf = svm.SVC(kernel='rbf')
svm_rbf.fit(X_train, y_train)
```

```
Out[18]: SVC()
```

```
In [19]: # predict
predictions = svm_rbf.predict(X_test)

# accuracy
print(metrics.accuracy_score(y_true=y_test, y_pred=predictions))
```

```
0.9396130952380952
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js